

Chapter 14: Basics of Functional Dependencies and Normalization for Relational Databases

Introduction

- So far we have assumed:
 - **Attributes** are grouped to form a relation schema by using the common sense of the database designer
- Need some **formal way** of analyzing why one grouping of attributes into a relation schema may be better than another
- In this chapter we discuss some of the **theory** to evaluate relational schemas for design quality

What Are Some Database Schema Better Than Others

- Consider the following a very simple way to **design** a **database schema**:
 - Put **every attribute** that you need to store **into one single (huge) relation**.

- Example: the University database schema will become:

```
Company (SSN, fname, ..., DNo, DName, ..., PNo, PName, ..., DepName,...)
```

- This relation is called the **Universal Relation**
- So what is **bad** about this relation?

What Makes a Relation “Good” or “Bad”?

- **Four informal guidelines:** **measures** to determine **the quality of** relation schema design:
 - Making sure that the **semantics of the attributes is clear** in the schema
 - Reducing the **redundant** information in tuples
 - Reducing the **NULL** values in tuples
 - Disallowing the possibility of generating **spurious tuples**

Imparting Clear Semantics to Attributes in Relations

- Making sure that the **semantics of the attributes is clear** in the schema
 - Design a relation schema so it is **easy to explain its meaning**
 - Do not **combine** attributes from **multiple entity types** and **relationship types** **into** a **single relation**.
 - If the relation corresponds to a **mixture of multiple entities**, **semantic ambiguities** will result, and the relation **cannot** be easily explained

Imparting Clear Semantics to Attributes in Relations (Example)

EMPLOYEE

F.K.

Ename	<u>Ssn</u>	Bdate	Address	Dnumber
-------	------------	-------	---------	---------

P.K.

PROJECT

F.K.

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

P.K.

EMP_DEPT

Ename	<u>Ssn</u>	Bdate	Address	Dnumber	Dname	Dmgr_ssn
-------	------------	-------	---------	---------	-------	----------

EMP_PROJ

<u>Ssn</u>	<u>Pnumber</u>	Hours	Ename	Pname	Plocation
------------	----------------	-------	-------	-------	-----------

Redundant Information in Tuples

EMP_DEPT

Ename	<u>Ssn</u>	Bdate	Address	Dnumber	Dname	Dmgr_ssn
Smith, John B.	123456789	1965-01-09	731 Fondren, Houston, TX	5	Research	333445555
Wong, Franklin T.	333445555	1955-12-08	638 Voss, Houston, TX	5	Research	333445555
Zelaya, Alicia J.	999887777	1968-07-19	3321 Castle, Spring, TX	4	Administration	987654321
Wallace, Jennifer S.	987654321	1941-06-20	291 Berry, Bellaire, TX	4	Administration	987654321
Narayan, Ramesh K.	666884444	1962-09-15	975 FireOak, Humble, TX	5	Research	333445555
English, Joyce A.	453453453	1972-07-31	5631 Rice, Houston, TX	5	Research	333445555
Jabbar, Ahmad V.	987987987	1969-03-29	980 Dallas, Houston, TX	4	Administration	987654321
Borg, James E.	888665555	1937-11-10				

Redundancy

Redundancy

Redundancy

EMP_PROJ

<u>Ssn</u>	<u>Pnumber</u>	Hours	Ename	Pname	Plocation
123456789	1	32.5	Smith, John B.	ProductX	Bellaire
123456789	2	7.5	Smith, John B.	ProductY	Sugarland
666884444	3	40.0	Narayan, Ramesh K.	ProductZ	Houston
453453453	1	20.0	English, Joyce A.	ProductX	Bellaire
453453453	2	20.0	English, Joyce A.	ProductY	Sugarland

Update Anomalies

- Storing such **joins of base relations** leads to an additional problem referred to as **update anomalies**
- Database researchers has found a number of **bad properties** that they call **anomalies**
 - When a relation exhibits one or more of these "anomalies", it is deemed to be "bad"
- There are three types of database anomalies:
 - Insert anomaly
 - Delete anomaly
 - Update anomaly

Insert Anomaly

- See example:

SSN	FName	LName	BDate	DNum	DName	MgrSSN
111-11-1111	John	Smith	Jan-1-78	5	Research	123-45-6789
222-22-2222	Jane	Doe	Apr-1-76	5	Research	123-45-6789
333-33-3333	Jack	Rabbit	May-4-79	1	Payroll	777-77-7777

- If we **insert a new tuple** for an **employee** who works in **department number 5**
 - Must enter **all** the attribute values of department 5 correctly so that they are *consistent* with other tuples

Insert Anomaly (cont.)

- See example:

				SSN	FName	LName	BDate	DNum	DName	MgrSSN
				111-11-1111	John	Smith	Jan-1-78	5	Research	123-45-6789
111-11-1111	John	Smith	Jan	222-22-2222	Jane	Doe	Apr-1-76	5	Research	123-45-6789
222-22-2222	Jane	Doe	Apr	333-33-3333	Jack	Rabbit	May-4-79	1	Payroll	777-77-7777
333-33-3333	Jack	Rabbit	Ma	79				6	HR	7777

- If we insert a department, say (dnumber=6, dname='Human Resources'), that **does not have any employees**, then we **need to use NULL values**

Delete Anomaly

- See example:

SSN	FName	LName	BDate	DNum	DName	MgrSSN
111-11-1111	John	Smith	Jan-1-78	5	Research	123-45-6789
222-22-2222	Jane	Doe	Apr-1-76	5	Research	123-45-6789
333-33-3333	Jack	Rabbit	May-4-79	1	Payroll	777-77-7777

- If we **delete** the **last employee** "Jack Rabbit", we will **lose information on the department**

DELETE employee

WHERE fname = 'Jack' AND lname = 'Rabbit'

SSN	FName	LName	BDate	DNum	DName	MgrSSN
111-11-1111	John	Smith	Jan-1-78	5	Research	123-45-6789
222-22-2222	Jane	Doe	Apr-1-76	5	Research	123-45-6789

Delete Anomaly (cont.)

- This delete operation has deleted **additional** information !!!
 - Information on the **"Payroll" department** is **also deleted !!!** (its dnumber, its name and its manager) from the database !!!
 - Clearly, that should **not** be a consequence of a delete command that is routinely used to remove employees....

Update Anomaly

- See example:

		SSN	FName	LName	BDate	DNum	DName	MgrSSN
		111-11-1111	John	Smith	Jan-1-78	5	Research	888-88-8888
SSN	FName	222-22-2222	Jane	Doe	Apr-1-76	5	Research	888-88-8888
111-11-1111	John	333-33-3333	Jack	Rabbit	May-4-79	1	Payroll	777-77-7777
222-22-2222	Jane	Doe	Apr-1-76	5	Research	123-45-6789		
333-33-3333	Jack	Rabbit	May-4-79	1	Payroll	777-77-7777		

- If we change the manager of the Research department to 888-88-8888
- The update **logically** involved changing ONE item of information (namely change 123-45-6789 to 888-88-8888), **but** the update operation **has modified multiple tuples** in Employee

Constraints Prevent (some) Anomalies in the Data

Student	Course
Mary	CS145
Joe	CS145
Sam	CS145
..	..

Course	Room
CS145	B01
CS229	C12

Is this form better?

- Redundancy?
- Update anomaly?
- Delete anomaly?
- Insert anomaly?

Today: develop theory to understand why this design may be better **and** how to find this *decomposition*...

Design Guideline 1

- Design guideline 1:
 - Design the base relation schemas so that **no insertion, deletion, or modification anomalies** are present in the relations.
 - **If any anomalies are present, note them** clearly and make sure that the programs that **update the database** will **operate correctly**

Design Guideline 2

- Design guideline 2:
 - As far as possible, **avoid** placing attributes in a base relation whose values may **frequently be NULL**.
 - **If** NULLs are unavoidable, make sure that they apply in **exceptional cases only** and **do not** apply to a **majority of tuples** in the relation

Design Guideline 3

- Design guideline 3:
 - Avoid relations that contain **matching attributes** that are **not (foreign key, primary key) combinations** because joining on such attributes may **produce spurious tuples**.

Summary and Discussion of Design Guidelines

- We have **informally** discussed situations that lead to **problematic** relation schemas
- We proposed **informal guidelines** for a good relational design
- In the rest: Present **formal** concepts that may be used to define the *goodness* and *badness* of relation schemas **more precisely**.
 - Discuss **Functional dependency** as a tool for analysis
 - Specify the three **normal forms**

Introduction to Functional Dependency

- A **formal tool** for **analysis** of **relational schemas**
 - Enables us to *detect* and *describe* some of the above-mentioned **problems** in **precise terms**

Functional Dependencies

- Let X and Y be 2 **set of attributes** in a relation R
- We say that Y is **functionally dependent on X** (or X **functionally determines Y** , notation: $X \rightarrow Y$) iff:

$X \rightarrow Y \iff$ for any two tuples $t1$ and $t2$ of the relation R ,
if $t1[X] = t2[X]$ (i.e., the attribute values for
the X attributes are same in both tuples)
then $t1[Y] = t2[Y]$

- See example:

Functional Dependencies (cont.)

- Example: A case where Y is *not* functionally dependent on X

Relation:			
X			
A	B	C	
...	
tuple 1:	...	b7	c4

tuple 2:	...	b7	c4

Relation:							
X				Y			
A	B	C	D	E	F	G	
...	
tuple 1:	...	b7	c4	...	e1	f3	g9

tuple 2:	...	b7	c4	...	e1	f3	g9

- What you *must see* if Y is functionally dependent on X

Example of Functional Dependencies

- Consider the following relation:
 - That represent information about the **employees** AND the **projects** that **they work on**

Employee1(SSN, FName, LName, PNumber, PName, Hours)

The key of this relation is: (SSN, PNumber)

- Sample content of the Employee1 relation:

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

Example of Functional Dependencies (cont.)

- Some **functional dependencies** in the Employee1 relation:
 - $SSN \rightarrow \text{fname, lname}$
 - $PNumber \rightarrow PName$
 - $SSN, PNumber \rightarrow \text{Hours}$
 - But also: $SSN, PNumber \rightarrow \text{fname, lname} !!!$
 - And also: $SSN, PNumber \rightarrow PName !!!$

Normal Forms

- How to use **functional dependencies** to **develop a formal methodology** for **testing** and **improving** relation schemas
- Takes a relation schema through **a series of tests** to ***certify*** **whether it satisfies a certain normal form**
 - 1NF, 2NF and 3NF, which are based on primary keys

Normal Forms

- 1st Normal Form (1NF) = All tables are flat

- 2nd Normal Form (2NF)

- Boyce-Codd Normal Form (BCNF)

- 3rd Normal Form (3NF)

DB designs based on
functional dependencies,
intended to prevent
data ***anomalies***

- 4th and 5th Normal Forms = see text books

The First Normal Form (1NF)

- The 1NF is the **simplest** and the **only one** that **does not** depend on the notion of "**functional dependency**"
- A relation is in 1NF (first normal form) if:
 - Every attribute of the relation has atomic (single, not multi) values

1st Normal Form (1NF)

Student	Courses
Mary	{CS145,CS229}
Joe	{CS145,CS106}
...	...

Violates 1NF.

Student	Courses
Mary	CS145
Mary	CS229
Joe	CS145
Joe	CS106

In 1st NF

1NF Constraint: Types must be atomic!

Full Functional Dependency

- $X \rightarrow Y$ is a **full functional dependency** if **removal** of any attribute A **from X** means that the **dependency does not hold anymore**
- That is, for any attribute $A \in X$, $(X - \{A\})$ **does not** functionally determine Y .
- Example: $\{\text{Ssn}, \text{Pnumber}\} \rightarrow \text{Hours}$ is a full dependency
 $\{\text{Ssn}, \text{Pnumber}\} \rightarrow \text{Ename}$ is partial dependency
- 2NF is based on the concept of **full functional dependency**

The Second Normal Form (2NF)

- A relation is in **2NF** iff:
 - **Relation** is in **1NF** (i.e., every attribute is atomic), and
 - **Every non-key attribute** is **fully functionally determined** by **every key** of the relation


$\forall N \notin \text{KEY}: K \rightarrow N \text{ and } K-A \text{ does not } \rightarrow N \text{ (K is a key)}$

- In other words

- You must **not** have a **non-key** attribute that is **functionally determined** by **only a (proper) subset** of a key

The Second Normal Form (2NF) (cont.)


- How to **spot** the fact that a **relation violates** the **2NF**
 - Find a **non-key attribute** that is **functionally determined** by a **subset of some key**
 - Graphically:



This will make the relation **violates the 2NF** criteria.

Meaning of the 2NF

- The 2NF deals with **embedded entities**:
 - A **relation** that **violates** the **2NF** contains another embedded autonomous entity



There may be **other attributes** (besides F) that are **functionally dependent** on **B** !!!

Example of a Violation of the 2NF

- Consider the following relation:

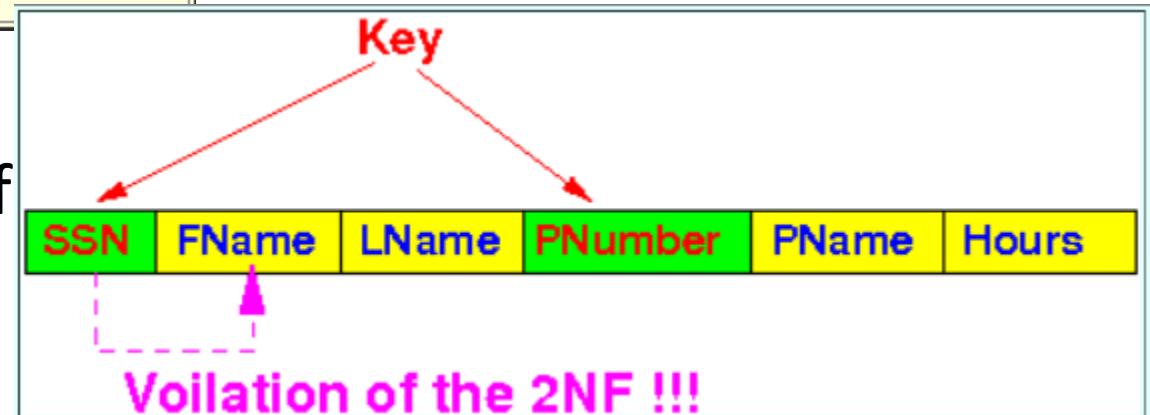
Employee1(SSN, FName, LName, PNumber, PName, Hours)

Functional Dependencies:

SSN → FName, LName
PNumber → PName
SSN, PNumber → Hours

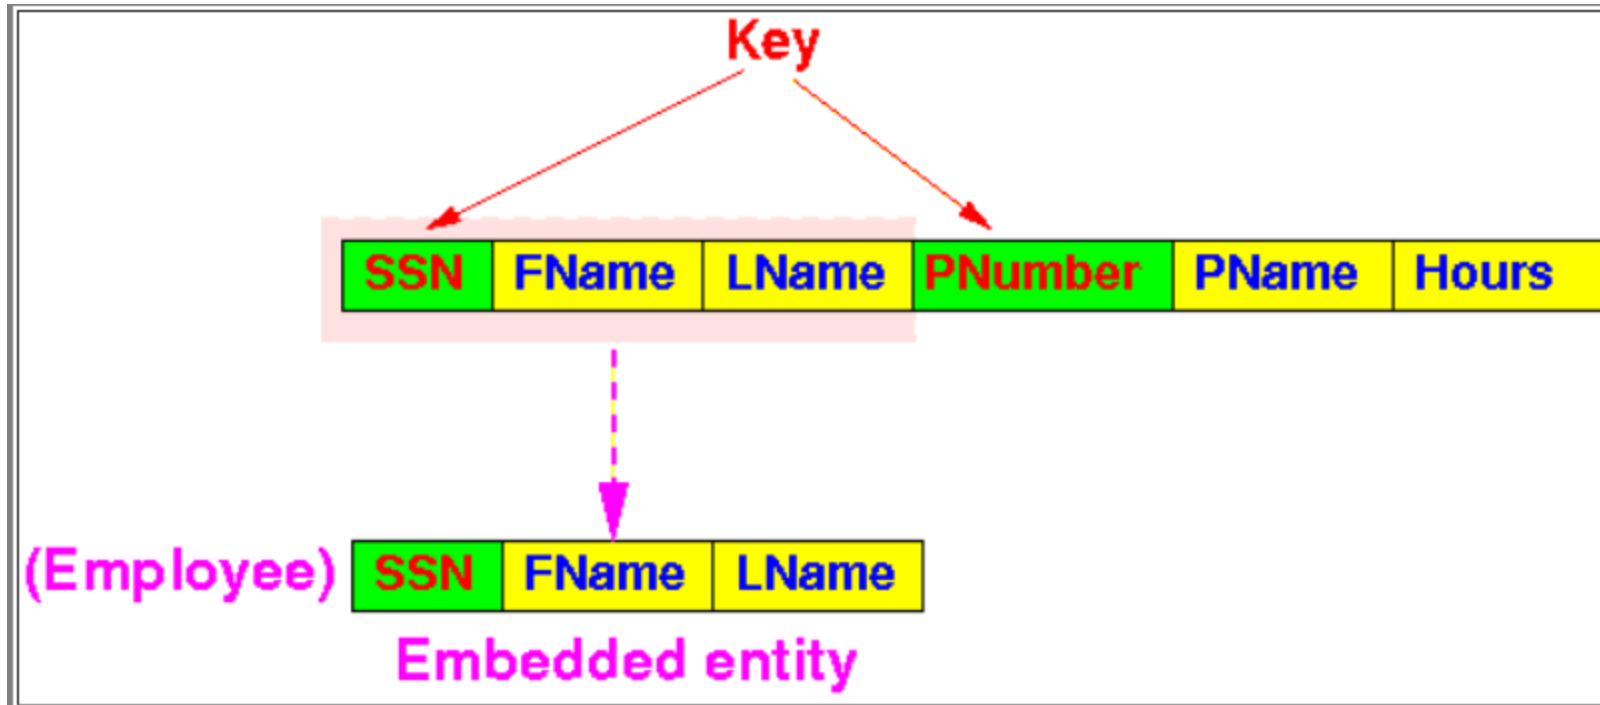
Keys: (SSN, PNumber)

- Employee1 **is not** in **2NF** because of



Example of a Violation of the 2NF (cont.)

- The **entity** that is **embedded** inside this **relation** is **Employee**:



How to Solve Normal Form Violations : Decomposition

- **How** can we **remove** the **violation** of a **normal form criteria**
- **Answer:** **Decompose** (= break up) a **relation** into **two or more relations**
- Example:

Employee1(SSN, FName, LName, PNumber, PName, Hours)

SSN → FName causes 2NF violation

A decomposition:

R1(PNumber, PName, Hours)

R2(SSN, LName, FName)

Good and Bad Decomposition

- We need to **learn** more about the **effect** of **breaking a relation** into **2 or more relations**
- Fact:
 - There are **good** and **bad decomposition**
- Note
 - The **decomposition** suggested by the previous example was a **bad** decomposition

Decomposition and Its Effect on a Relation

- A **decomposition** of a relation **R** is :
 - There are **good** and **bad decomposition**

▪ **Decomposition** = a **collection** of relations **R_1, R_2, \dots, R_n** , such that:

1. **Each attribute** in **R_1, R_2, \dots, R_n** is **found** in **R** and

2. **Every attribute** of **R** appears in **R_1, R_2, \dots, R_n** **at least once**

- In other words:
 - A **decomposition** will **not lose** any attributes from the original relation

Decomposition and Its Effect on a Relation (cont.)

- There are "good" and "bad" decompositions
- We will soon see that

- A "bad" decompositions will *introduce extraneous information* into the **database** !!!
- A "good" decompositions will *preserve* the **content** of the *original relation*

The Effect of a Decomposition

- Consider the following content of the **relation** Employee1:

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

- The content of the relation Employee1 convey the following **facts**:

- There are 2 employees

- (111-11-1111, John, Smith)
- (111-22-3333, Jane, Doe)

- There are 2 projects

- (pj1, DBApplet)
- (pj2, WebServer)

- There are 3 "works_on" information items:

- (111-11-1111, pj1, 20),
- (111-11-1111, pj2, 10)
- (111-22-3333, pj1, 5)

The Effect of a Decomposition (cont.)

- Go back to our previous example:

A decomposition for **Employee1**(SSN, FName, LName, PNumber, PName, Hours)

R1(SSN, LName, FName)

R2(PNumber, PName, Hours)

- Now **populate R1** and **R2** using the data in **Employee1**:

$R1 = \pi_{SSN, FName, LName}(Employee1)$

$R2 = \pi_{PNumber, PName, Hours}(Employee1)$

R1

SSN	FName	LName
111-11-1111	John	Smith
111-22-3333	Jane	Doe

R2

PNumber	PName	Hours
pj1	DBApplet	20
pj2	WebServer	10
pj1	DBApplet	5

- Therefore **R2** will be:

The Effect of a Decomposition (cont.)

- Note:
 - After the decomposition of **Employee1** into **R1** and **R2**, the relation **Employee1** is deleted !!!
- An important question to ask is:
 - Can we obtain the same information stored in **Employee1** from the relations **R1** and **R2** ?
 - (If not, we will be in deep trouble)...

Reconstructing the Original Content of a Decomposed Relation

- The **reconstruction algorithm** used is as follows:

```
if (  $R1 \cap R2 \neq \emptyset$  )  
{  
    reconstruction =  $R1 \bowtie R2$  // Join  
}  
else  
{  
    reconstruction =  $R1 \times R2$  // Cartesian product  
}
```

Reconstructing the Original Content of a Decomposed Relation (Example)

R1			R2		
SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-22-3333	Jane	Doe		WebServer	10
			pj1	DBApplet	5

Reconstruction:

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-11-1111	John	Smith	pj1	DBApplet	5
111-22-3333	Jane	Doe	pj1	DBApplet	20
111-22-3333	Jane	Doe	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

- We performed a **cartesian product** $R1 \times R2$ on **R1** and **R2** because they do **not** have any attributes **in common** !

Reconstructing the Original Content of a Decomposed Relation (Example) (cont.)

- Compare the **content** of the **reconstruction** to the **content** of the **original** Employee1 relation:

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

Reconstruction:

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-11-1111	John	Smith	pj1	DBApplet	5
111-22-3333	Jane	Doe	pj1	DBApplet	20
111-22-3333	Jane	Doe	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

- The **good**:

- We are able to **obtain *every tuple*** that was in the **original relation** :

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

- The **bad**:

- There are ***extraneous tuples*** that were **not present** in the ***original relation***

Reconstruction: (**extraneous** tuples in **red**)

SSN	FName	LName	PNumber	PName	Hours
111-11-1111	John	Smith	pj1	DBApplet	20
111-11-1111	John	Smith	pj2	WebServer	10
111-11-1111	John	Smith	pj1	DBApplet	5
111-22-3333	Jane	Doe	pj1	DBApplet	20
111-22-3333	Jane	Doe	pj2	WebServer	10
111-22-3333	Jane	Doe	pj1	DBApplet	5

Requirements of a decomposition

- After **decomposing** a **relation R**:
 - We **must** be able to **obtain all tuples** in the **original relation R** using the reconstruction algorithm
 - We **must not** obtain **extraneous tuples** that were **not** present in the **original relation R** using the reconstruction algorithm

Requirements of a decomposition (cont.)

- If we **miss** some **tuples** in the reconstruction, it means that:
 - We have **lose information**
 - **Clearly** that is **unacceptable!!!**
- If we **gain** some (**extraneous**) **tuple** in the reconstruction, it means that:
 - We have some **invalid information** in the relation (= database) !!!
 - That is **also unacceptable !!!**